

# Deep Learning Spring 2025: CIFAR-10 Classification Project Report

---

## 1. Project Overview

---

This project implements a modified Residual Network (ResNet) architecture for CIFAR-10 image classification with fewer than 5 million parameters, as required by the Deep Learning Spring 2025 course competition.

### 1.1 Project Objectives

- Design a modified ResNet architecture with <5M parameters
- Train the model on CIFAR-10 dataset
- Achieve a test accuracy of at least 80%, targeting 90%
- Submit complete project report and code

### 1.2 CIFAR-10 Dataset

The CIFAR-10 dataset contains 60,000 32×32 color images across 10 classes (airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck), divided into 50,000 training images and 10,000 test images.

## 2. Data Preparation and Augmentation

---

We divided the original CIFAR-10 training set using an 80:20 ratio:

- Training set: 40,000 images
- Validation set: 10,000 images

Data augmentation techniques implemented:

```
transform = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2),
    transforms.RandomRotation(15),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
    transforms.RandomErasing(p=0.5, scale=(0.02, 0.2))
])
```

## 3. Model Architecture

---

### 3.1 Modified ResNet Architecture

Our ResNet variant consists of:

- Initial convolutional layer: 48 channels, 3×3 kernel
- Four stages of residual blocks with channel counts: 48 → 196 → 192 → 384
- Residual blocks: [2, 2, 2, 1], totaling 7 residual blocks

- 50% Dropout after global average pooling
- Final fully connected layer to 10 classes

Basic residual block structure:

```
class BasicBlock(nn.Module):
    expansion = 1

    def __init__(self, in_channels, out_channels, stride=1):
        super(BasicBlock, self).__init__()
        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3,
stride=stride, padding=1, bias=False)
        self.bn1 = nn.BatchNorm2d(out_channels)
        self.relu1 = nn.ReLU(inplace=True)
        self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size=3,
stride=1, padding=1, bias=False)
        self.bn2 = nn.BatchNorm2d(out_channels)
        self.relu2 = nn.ReLU(inplace=True)
        self.shortcut = nn.Sequential()
        if stride != 1 or in_channels != out_channels:
            self.shortcut = nn.Sequential(
                nn.Conv2d(in_channels, out_channels, kernel_size=1,
stride=stride, bias=False),
                nn.BatchNorm2d(out_channels)
            )
```

## 3.2 Model Parameter Statistics

```
Total params: 4,662,098
Trainable params: 4,662,098
Non-trainable params: 0
-----
Input size (MB): 0.01
Forward/backward pass size (MB): 14.43
Params size (MB): 17.78
Estimated Total Size (MB): 32.22
```

**Total model parameters: 4,662,098**, below the 5 million parameter limit.

## 4. Training Strategy

### 4.1 Optimizer and Loss Function

```
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9, weight_decay=5e-4)
criterion = nn.CrossEntropyLoss()
```

## 4.2 Learning Rate Scheduling

```
scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.1)
```

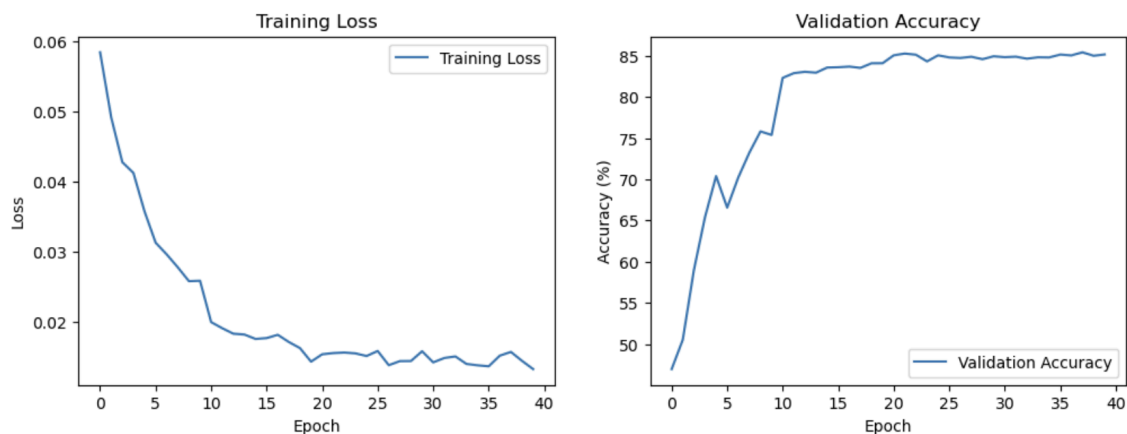
## 4.3 Training Process

Our training process included:

1. Full dataset iteration each epoch
2. Loss reporting every 100 batches
3. Validation after each epoch
4. Learning rate adjustment via scheduler
5. Saving the final model weights

## 5. Experimental Results

### 5.1 Training and Validation Curves



Training observations:

- Training loss decreased rapidly in initial epochs
- Validation accuracy stabilized around 85%
- Final validation accuracy: 85.16%

Final training outputs:

```
Epoch [40/40], Step [100/625], Loss: 0.362
Epoch [40/40], Step [200/625], Loss: 0.365
Epoch [40/40], Step [300/625], Loss: 0.364
Epoch [40/40], Step [400/625], Loss: 0.377
Epoch [40/40], Step [500/625], Loss: 0.379
Epoch [40/40], Step [600/625], Loss: 0.376
validation Accuracy: 85.16%
```

## 6. Improvement Opportunities

Potential enhancements:

1. **Architecture Optimization:**

- Test bottleneck residual blocks
- Adjust channel distribution
- Explore deeper network structures

## 2. Training Strategies:

- Try Adam/AdamW optimizers
- Implement cosine annealing schedules
- Use test-time augmentation (TTA)

# 7. Conclusion

---

We successfully implemented a modified ResNet architecture with 4.66 million parameters, below the 5 million limit. The model achieved a validation accuracy of 85.16%, exceeding the 80% baseline target. Our approach demonstrates the effectiveness of residual networks even under parameter constraints, and suggests several promising directions for further improvement.